

## Headlamp Multi-Cluster Setup Guide (MicroK8s + Headlamp)

This document explains how to integrate **multiple MicroK8s clusters** into **Headlamp**

### Overview

Headlamp is a modern, extensible Kubernetes UI that allows you to monitor and manage multiple Kubernetes clusters from a single interface.

### How Multi-Cluster Works

Headlamp connects to Kubernetes clusters using **kubeconfig files**. For **multi-cluster mode**, Headlamp reads a **single kubeconfig** that contains **multiple cluster entries**.

Flow:

1. Individual MicroK8s clusters generate kubeconfig files
2. Kubeconfigs are merged into a single multi-cluster config
3. Config is stored as a Kubernetes Secret
4. Headlamp reads the Secret and presents the cluster switcher UI

### Architecture

#### Component Layout

Component	Location	Purpose
Headlamp	Cluster A	Web UI & cluster management
Merged Kubeconfig	Secret in Cluster A	Multi-cluster authentication
API Servers	All clusters	Kubernetes control plane endpoints

Ingress Controller	Cluster A	HTTPS access to Headlamp UI
--------------------	-----------	-----------------------------

## Prerequisites

Component	Installation
MicroK8s	<b>sudo snap install microk8s --classic</b>
kubectl	<b>sudo snap install kubectl --classic</b>
Headlamp	<b>can download Helm or YAML Configuration</b>

## Setup Paths

Choose your deployment approach:

### Path A: Quick Test (Development)

**Use when:**

- Testing multi-cluster functionality
- Development/staging environments

**Characteristics:**

- Uses insecure-skip-tls-verify: true
- No certificate management
- Public IPs for API access
- NodePort for UI access
- Not production-ready

### Path B: Production Secure

**Use when:**

- Production deployments

- Security compliance required
- Long-term installations

**Characteristics:**

- Full TLS certificate validation
- VPC peering or private network
- Ingress with Let's Encrypt
- RBAC and authentication
- Production-ready

## Installation Steps

### Step 1: Choose Headlamp Host Cluster

Decision: Which cluster will run Headlamp?

```
# On Cluster A (chosen as Headlamp host)
hostname
# Example: cluster-a-server
# Verify MicroK8s is running
microk8s status --wait-ready
```

Important: All subsequent Headlamp operations (merge, secret creation, deployment) happen on this cluster.

### Step 2: Configure API Server Accessibility

Goal: Make each cluster's API server reachable from the Headlamp pod.

#### 2.1 Determine Your Network Strategy

- Option A: Private IPs with VPC Peering (Production)  
 Option B: Public IPs (Development/Quick Test)

#### 2.2 Configure API Server (Run on EACH cluster)

```
# Stop MicroK8s
sudo microk8s stop
# Edit API server arguments
```

```
sudo nano /var/snap/microk8s/current/args/kube-apiserver
```

**Add or modify these lines:**

```
# For all interfaces (development)
--advertise-address=0.0.0.0
# OR for specific IP (production)
--advertise-address=172.31.7.94
# Ensure secure port is set
--secure-port=16443
```

### **2.3 Configure Cloud Firewall/Security Groups**

**(run for EACH cluster):**

Navigate to EC2 → Security Groups → Select your cluster's SG

**Add Inbound Rules:**

Type	Protocol	Port	Source	Description
Custom TCP	TCP	16443	<b>Cluster A IP/32 or 0.0.0.0/0 for testing</b>	<b>Allow Headlamp to access API</b>
Custom TCP	TCP	16443	<b>Cluster B IP/32 or 0.0.0.0/0 for testing</b>	<b>Allow inter-cluster access</b>

### **2.4 Verify API Accessibility**

```
# From Cluster A, test reaching Cluster B
curl -k https://<cluster-b-ip>:16443/version
# Expected: JSON response with version info
# Error: Connection refused = firewall issue
# Error: Connection timeout = routing issue
```

### **Step 3: TLS Strategy Selection**

Choose your TLS approach:

#### **Option A: Secure TLS (Production)**

Pros:

Full certificate validation

Production-ready

Cons:

More complex setup

Requires certificate regeneration

When to use: Production, compliance-required environments

#### **Option B: Insecure Skip TLS (Development)**

Pros:

Faster setup

No certificate management

Cons:

Man-in-the-middle vulnerability

Not production-ready

Security audit failures

When to use: Development, testing, PoCs only

### **Step 4: Configure TLS Certificates (Option A Only)**

Skip this step if using Option B

#### **4.1 Add Subject Alternative Names (SANs)**

**Run on EACH cluster:**

# Edit certificate template

```
sudo nano /var/snap/microk8s/current/certs/csr.conf.template
```

**Under [ alt\_names ] section, add:**

```
[alt_names]
DNS.1 = kubernetes
DNS.2 = kubernetes.default
DNS.3 = kubernetes.default.svc
DNS.4 = kubernetes.default.svc.cluster.local
DNS.5 = cluster-a.yourdomain.com
# Add your DNS
IP.1 = 127.0.0.1
IP.2 = 172.31.7.94# Private IP
IP.3 = 13.239.17.107# Public IP (if used)
```

## 4.2 Regenerate Certificates

```
# Backup existing certificates
sudo cp -r /var/snap/microk8s/current/certs /var/snap/microk8s/current/certs.backup
# Stop MicroK8s
sudo microk8s stop
# Remove old certificates
sudo rm /var/snap/microk8s/current/certs/server.crt
sudo rm /var/snap/microk8s/current/certs/server.key
# Start MicroK8s (automatically regenerates certs)
sudo microk8s start
microk8s status --wait-ready
# Verify new certificate includes SANs
openssl x509 -in /var/snap/microk8s/current/certs/server.crt -text -noout | grep -A 10
"Subject Alternative Name"
```

## Step 5: Export Kubeconfig from Each Cluster

### 5.1 Export from Cluster A

```
# Export kubeconfig
```

```
microk8s config > ~/cluster-a.yaml  
# Verify export  
cat ~/cluster-a.yaml | grep "server:"  
# Should show: server: https://127.0.0.1:16443
```

## 5.2 Export from Cluster B

```
# Export kubeconfig  
microk8s config > ~/cluster-b.yaml  
# Verify export  
cat ~/cluster-b.yaml | grep "server:"
```

## 5.3 Transfer Configs to Cluster A

```
# From Cluster A, copy Cluster B's config  
scp ubuntu@<cluster-b-ip>:~/cluster-b.yaml ~/cluster-b.yaml  
# Verify both files exist on Cluster A  
ls -lh ~/cluster-*.yaml
```

## Step 6: Edit Kubeconfig Files

Make each kubeconfig unique and accessible.

### 6.1 Edit cluster-a.yaml

```
nano ~/cluster-a.yaml  
apiVersion: v1  
kind: Config  
clusters:  
- cluster:  
  certificate-authority-data: LS0tLS1CRU... # Keep existing  
  # For Option B (insecure), add:  
  insecure-skip-tls-verify: true  
  # Change server from 127.0.0.1 to accessible IP:  
  server: https://x.x.x.x:16443 # Public IP or DNS
```

```

name: microk8s-cluster-a # Make unique

contexts:
- context:
  cluster: microk8s-cluster-a # Match cluster name
  user: admin-cluster-a # Make unique
  name: cluster-a # Make unique
  current-context: cluster-a

users:
- name: admin-cluster-a # Make unique
user:
client-certificate-data: LS0tLS1... # Keep existing
client-key-data: LS0tLS1... # Keep existing

```

## 6.2 Edit cluster-b.yaml

Repeat the same process with these unique names:

Cluster name: microk8s-cluster-b  
 Context name: cluster-b  
 User name: admin-cluster-b  
 Server: https://x.x.x.x:16443

## Step 7: Install System kubectl on Cluster A

MicroK8s kubectl ( microk8s kubectl ) doesn't read ~/.kube/config by default.  
 Install the system kubectl for kubeconfig management.

```

# Install system kubectl
sudo snap install kubectl --classic
# Remove any aliases that override kubectl
nano ~/.bashrc
# Remove line: alias kubectl='microk8s kubectl'
# Save and reload
source ~/.bashrc
# Verify you're using system kubectl
which kubectl

```

```
# Should show: /snap/bin/kubectl (not microk8s)
```

## Step 8: Merge Kubeconfigs

### 8.1 Create .kube Directory

```
mkdir -p ~/.kube
```

### 8.2 Move Config Files

```
# Move to standard location
```

```
mv ~/cluster-a.yaml ~/.kube/cluster-a.yaml
```

```
mv ~/cluster-b.yaml ~/.kube/cluster-b.yaml
```

### 8.3 Merge Configurations

```
Using KUBECONFIG environment variable
```

```
export KUBECONFIG=~/.kube/cluster-a.yaml:~/.kube/cluster-b.yaml
```

```
kubectl config view --merge --flatten > ~/.kube/multi-cluster.yaml
```

### # Unset environment variable

```
unset KUBECONFIG
```

### 8.4 Verify Merge

```
# View merged config
```

```
kubectl --kubeconfig=~/.kube/multi-cluster.yaml config view
```

```
# List clusters
```

```
kubectl --kubeconfig=~/.kube/multi-cluster.yaml config get-clusters
```

```
# Expected output:
```

```
# NAME
```

```
# microk8s-cluster-a
```

```
# microk8s-cluster-b
```

```
# List contexts
```

```
kubectl --kubeconfig=~/.kube/multi-cluster.yaml config get-contexts
```

## 8.5 Test Connectivity

```
# Test Cluster A
kubectl --kubeconfig=~/kube/multi-cluster.yaml --context=cluster-a get nodes
# Test Cluster B
kubectl --kubeconfig=~/kube/multi-cluster.yaml --context=cluster-b get nodes
# Both should return node lists
```

## Step 9: Create Headlamp Namespace

```
# Use system kubectl with merged config
cp ~/kube/multi-cluster.yaml ~/kube/config
# Switch to Cluster A context
kubectl config use-context cluster-a
# Verify you're on Cluster A
kubectl get nodes
# Create dedicated namespace for Headlamp
kubectl create namespace headlamp
# Verify namespace
kubectl get namespace headlamp
```

## Why a dedicated namespace?

- Separation of concerns
- Easier RBAC management
- Cleaner than mixing with kube-system

## Step 10: Create Kubeconfig Secret

```
# Create secret in headlamp namespace
kubectl create secret generic headlamp-kubeconfig \
--from-file=config=$HOME/.kube/multi-cluster.yaml \
-n headlamp
```

```
# Verify secret exists  
kubectl get secret headlamp-kubeconfig -n headlamp
```

## **Step 11: Deploy Headlamp**

**Option A: Using Helm**

**Option B: Using Manual YAML**

### **Edit the headlamp deployment - Option B**

```
args:  
- "-kubeconfig=/etc/headlamp/config"  
- "-plugins-dir=/headlamp/plugins"
```

```
volumeMounts:      # Add the volume mount  
- name: kubeconfig  
mountPath: /etc/headlamp
```

```
volumes:      # Add the volume  
- name: kubeconfig  
secret:  
secretName: headlamp-kubeconfig
```

### **# Apply the deployment**

```
kubectl apply -f ~/headlamp.yaml
```

## **Step 12: Verify Headlamp Deployment**

```
# Check pod status  
kubectl get pods -n headlamp
```

```
# Verify kubeconfig is mounted  
kubectl exec -n headlamp deployment/headlamp -- ls -la /etc/headlamp/
```

```
# Test if pod can access both clusters  
kubectl exec -n headlamp deployment/headlamp -- \  
kubectl --kubeconfig=/etc/headlamp/config config get-contexts
```

```
# Expected: Shows both cluster-a and cluster-b
```

### **Step 13: Configure Ingress with TLS**

For deployments, use Ingress with HTTPS:

```
# Enable MicroK8s ingress
microk8s enable ingress
# Enable cert-manager for Let's Encrypt
microk8s enable cert-manager
```

#### **# Create ClusterIssuer**

```
certificate-issuer.yaml
apiVersion: cert-manager.io/v1
kind: ClusterIssuer
metadata:
  name: lets-encrypt
spec:
  acme:
    email: your-email@example.com
    server: https://acme-v02.api.letsencrypt.org/directory
    privateKeySecretRef:
      name: lets-encrypt-private-key
    # Add a single challenge solver, HTTP01 using nginx
    solvers:
    - http01:
        ingress:
          class: public
kubectl apply -f ~/certificate-issuer.yaml
```

#### **# Create Ingress**

```
ingress.yaml
kind: Ingress
apiVersion: networking.k8s.io/v1
metadata:
```

```
name: headlamp-ingress
namespace: headlamp
annotations:
  cert-manager.io/cluster-issuer: "lets-encrypt"
spec:
  tls:
    - secretName: headlamp-tls-secret
      hosts:
        - headlamp.yourdomain.com
  rules:
    - host: headlamp.yourdomain.com
      http:
        paths:
          - path: /
            pathType: Prefix
        backend:
          service:
            name: headlamp
            port:
              number: 80
```

kubectl apply -f ~/ingress.yaml  
Access via: <https://headlamp.yourdomain.com>